

ID2223 – Lab 2, 2017

Deep Learning with TensorFlow

Introduction

Goals

This lab has the following goals:

- Learn how to setup and run a computational graph in Tensorflow
- Implement a single-layer as well as a multi-layer Neural Network in Tensorflow
- Combine different activation functions to increase the accuracy
- Tackle overfitting using regularization
- Further improve the performance by using Convolutional Layers
- Use hyperparameter optimization to improve prediction accuracy

Requirements

For this lab assignment in *Tensorflow*, we will use the [Tensorflow Python API](#). The version of Python you have installed on your machine must be at least 2.7. You can install Tensorflow [directly](#) on your machine, but we recommend to use the [Docker based installation](#) as it is less invasive (less things to install locally) and easier to make sure it works on every platform. For the docker based installation you will need to [install Docker Engine](#) first. The Docker image for Tensorflow contains [Jupyter IPython](#) which is an interactive environment Python. (Note: if you do not have a laptop, you will be able to use Hopsworks (www.hops.site) that supports the Jupyter/python with Tensorflow).

Fashion MNIST Dataset

For this lab we will use a dataset called *Fashion-MNIST*. The original MNIST dataset (Mixed National Institute of Standards and Technology database) is a database of hand-written digits that is commonly used for evaluating image classification algorithms. You can read more about the dataset in [Yann LeCun's MNIST page](#) or [Chris Olah's visualizations of MNIST](#). MNIST is considered the HelloWorld dataset for Deep Learning, however, for this lab, MNIST is too easy to get very high accuracy on.

As such, we will use Fashion-MNIST - a drop-in replacement for the original MNIST, released by Zalando. It contains images of various articles of clothing and accessories: shirts, bags, shoes, and other fashion items. The Fashion MNIST training set contains 55,000 examples, and the test set contains 10,000 examples. Each example is a 28x28 grayscale image (just like the images in the original MNIST), associated with a label from 10 classes (t-shirts, trousers, pullovers, dresses, coats, sandals, shirts, sneakers, bags, and ankle boots). You can read more about Fashion-MNIST in [Fashion-MNIST introduction](#) and in [GitHub](#).

If you work with Docker/IPython, the `template.py` program will download the training and test datasets (images and labels), creating a local `input/data` folder for the data. Alternatively, if you are using `www.hops.site`, you can follow the "TensorFlow tour" that will create the Fashion-MNIST dataset in the "`TestJob/data/mnist`" folder (stored in `tfrecord` format).

Tensorflow

The main building blocks of a Tensorflow program are: *Tensors*, a *Computational Graph*, *Variables*, *Constants*, *PlaceHolders*, *Fetches* and *Feeds*. These building blocks are well presented in the [Tensorflow's Documentation](#). Please read up on them and make sure that you understand how a typical Tensorflow program is written. Specifically, for Section 1 of this lab look at [Tensorflow Tutorial](#) and [Convolutional Netowrk](#) is useful for Section 6.

Useful Python Pointers

In this section, we give you some practical python pointers that are helpful for this lab. First, [this Python Tutorial](#) is a good baseline to learn the language of Python (e.g., have a quick look into section 4 and 5). [NumPy](#) is the fundamental package for scientific computing with Python. [NumPy Arrays](#) are the backbone of Tensors in Tensorflow (read about `ndarray` and broadcasting).

Run Tensorflow with Docker

Running the default starting command:

```
1 docker run -it -p 8888:8888 gcr.io/tensorflow/tensorflow
```

will not save your notebooks and you might lose your work. You should use named containers in order to save your work. For this, the first time you start the docker image, name it `tensorflow-id2223` by running:

```
1 docker run -p 8888:8888 --name tensorflow-id2223 -it gcr.io/tensorflow/↵
  ↵ tensorflow
```

And the next time you want to continue work on your docker image, run:

```
1 docker start -ai tensorflow-id2223
```

Code Template for Docker

The code that is used in all of the subsequent sections, contains a common sequence of steps. These common steps are presented in Listing 7. According to the requirements of each task, you should only complete the template code (numbered steps in the template) in order to complete the implementation. The learning and visualization steps are also given but not in this document, you can find them in our [github repo](#).

Listing 1: "Code Template"

```
# all tensorflow api is accessible through this
import tensorflow as tf

import matplotlib.pyplot as plt

from tensorflow.examples.tutorials.mnist import input_data

# load data
mnist = input_data.read_data_sets('input/data', one_hot=True)

# 1. Define Variables and Placeholders
# 2. Define the model
# 3. Define the loss function
# 4. Define the accuracy
# 5. Define an optimizer

# initialize
init = tf.initialize_all_variables()
sess = tf.Session()
sess.run(init)

# 6. Train and test the model, store the accuracy and loss per iteration
# 7a. If you are using Python/Docker, plot and visualise the accuracy and loss
# 7b. If you are using hops.site, write to Tensorboard logs, and visualize using Tensorboard
```

Code Template for Hops

If you decide to use Hops, you will not be running a pure python application, rather a PySpark application that will launch a TensorFlow application in a mapper function (fun() in the code listing below). The basic code for launching the TensorFlow application is as follows:

Listing 2: "Hops Code Template"

```
def fun():

    import tensorflow as tf
    from hops import tensorboard
    from hops import hdfs

    from tensorflow.examples.tutorials.mnist import input_data

    # load data
    mnist = input_data.read_data_sets('input/data', one_hot=True)

    # 1. Define Variables and Placeholders
    # 2. Define the model
    # 3. Define the loss function
    # 4. Define the accuracy
```

```

# 5. Define an optimizer

# initialize
init = tf.initialize_all_variables()
sess = tf.Session()
sess.run(init)

# 6. Train and test the model, store the accuracy and loss per iteration

# 7a. If you are using Python/Docker, plot and visualise the accuracy and loss
# 7b. If you are using hops.site, write to Tensorboard logs, and visualize using Tensorboard

# This code launches the TensorFlow application in a Spark Mapper
from hops import tflauncher
from hops import util
tb_hdfs_dir = tflauncher.launch(spark, fun)

```

1 One-Layer Softmax Regression

In this task, you are going to design a neural network with 1 layer of 10 softmax neurons. If you are new to python and Tensorflow, it might be useful to look at the [beginner part](#) of the Tensorflow Tutorial.

You should build your model based on the softmax regression formula:

$$Y = \text{softmax}(X * W + b)$$

We will use cross entropy as the loss (error) for evaluating the model.

$$H(Y', Y) = - \sum Y'_i \cdot \log Y_i$$

Due to the fact that you train in batches, you will compute the mean of the cross entropy over the trained batches.

Finally you should compute the correctness of your prediction for each image as a 0/1 result and then compute the accuracy as the mean over the predictions.

When you compute accuracy and loss you run the whole training/testing dataset through your current model and this is quite expensive. Doing so every iteration would make it impractical, so you will compute it every 100 rounds, which we can call epochs. Plot accuracy and loss for each epoch to see how your model evolves during the training iterations

The images from the Fashion-MNIST dataset are defined as 3 dimensional tensors : (*image_height*, *image_width*, *input_channel*) with the parameters:

- image shape 28 X 28
- input channels 1 - gray-scale images.

The softmax function that you use for each of the neurons takes as input a vector of values. Since our images are defined as 3 dimensional tensors you will need to reshape it into a vector with *image_height* * *image_width* * *input_channel* elements.

Listing 3: "One-layer softmax regression"

```
# 1. Define Variables and Placeholders
#the first dimension (None) will index the images in the mini-batch
X = tf.placeholder(tf.float32, [None, ?, ?, ?])
Y_ = ? # correct answers(labels)
W = tf.Variable(tf.zeros([784, 10])) # weights W[784, 10] 784=28*28
b = ? # biases b[10]
XX = tf.reshape(X, [-1, 784]) # flatten the images into a single line of pixels

# 2. Define the model - compute predictions
Y = tf.nn.softmax(?)

# 3. Define the loss function
cross_entropy = ?

# 4. Define the accuracy
accuracy = ?

# 5. Train with an Optimizer
train_step = tf.train.GradientDescentOptimizer(<learning_rate>).minimize(<loss_function>)
```

Task 1

As per listing 3 perform the programming tasks:

- Define the variables and placeholders
- Define a model
- Define the loss
- Define the accuracy
- Train with the GradientDescentOptimizer and a learning rate of 0.5
- Train with the AdamOptimizer (a slightly better optimizer) and a learning rate of 0.005

What loss and accuracy do you get when training over 10.000 iterations?

2 Feedforward Neural Network with Back-Propagation

Now we are going to improve the accuracy by adding 4 more layers into our Neural Network. In a multi-layer, feedforward Neural Network, each layer instead of doing a weighted sum of all pixels, it does a weighted sum of the output of its previous layer. Design your layers such that you will have 200, 100, 60, 30 and 10 neurons for each layer respectively.

You can choose a different activation function for each hidden layer. For this task we want you to investigate two different combinations of activation function: I) *sigmoid* and *softmax*, and II) *ReLU* and *softmax*. We always apply softmax in the last layer (the output layer) and the other activation functions in the hidden layers.

Now complete the code in Listing 4, use the 2 different settings for the activation function and learn your model. In the end, compare the convergence rate by looking at the accuracy and loss plot.

Listing 4: "5-layers with a separate activations"

```
# 1. Define Variables and Placeholders

X = tf.placeholder(tf.float32, [None, ?, ?, ?]) #the first dimension (None) will index the images
Y_ = ? # correct answers

# Weights initialised with small random values between -0.2 and +0.2
W1 = tf.Variable(tf.truncated_normal([784, ?], stddev=0.1)) # 784 = 28 * 28
B1 = tf.Variable(tf.zeros([?]))
W2 = ?
B2 = ?
W3 = ?
B3 = ?
W4 = ?
B4 = ?
W5 = ?
B5 = ?

# 2. Define the model
XX = ?
Y1 = tf.nn.? (tf.matmul(XX, W1) + B1)
Y2 = ?
Y3 = ?
Y4 = ?
Ylogits = ?
Y = tf.nn.? (Ylogits)

# 3. Define the loss function

cross_entropy = tf.nn.? (Ylogits, Y_) # calculate cross-entropy with logits
cross_entropy = tf.reduce_mean(?)
```

Task 2

Do the following tasks in your program:

1. initialize all of the weights and biases
2. redefine your model by connecting the output of each layer to the input of the next layer
3. calculate the Logits (scores)
4. use a specific cross-entropy function to calculate loss using logits

Finally, analyse your results and try to answer the following questions:

Questions

1. What is the maximum accuracy that you can get in each setting for running your model with 10000 iterations?
2. Is there a big difference between the convergence rate of the sigmoid and the ReLU? If yes, what is the reason for the difference?
3. What is the reason that we use the softmax in our output layer?
4. By zooming into the second half of the epochs in accuracy and loss plot, do you see any strange behaviour? What is the reason and how you can overcome them? (e.g., look at fluctuations or sudden loss increase after a period of decreasing loss).

3 Regularization and Tuning

In the result of the model in the previous section, you probably can see signs of overfitting. In this section, we want you to try the techniques such as *dropout* and *learning rate decay* to see whether you can improve the accuracy of your model. You are, of course, free to add other regularization techniques (such as L2 regularization).

Task 3

The programming tasks involve adding to your neural network:

1. learning rate decay
2. regularization, such as dropout
3. both

An example is given in listing 5 on how to tell Tensorflow to apply dropout to layer1. It defines Y1d and uses this as input for Y2, instead of the previous Y1. When you run it, the model will expect to receive a value for the pkeep placeholder. You will need to provide a value for it in the `sess.run`

When implementing learning rate decay, you need to tell the optimizer to use a placeholder for the learning rate as it will change during the iterations and it will be provided at runtime. You will thus need to provide a new learning rate value for the placeholder every time you call the `sess.run`. Alternatively you can define a variable for the learning rate using the `tf.train.exponential_decay`

Listing 5: "Dropout and learning rate decay"

```
# 1. Define Variables and Placeholders
#learning rate placeholder
lr = ?
# placeholder for probability of keeping a node during dropout = 1.0 at test time (no dropout) and 0.75 at ↵
  ↳ training time
pkeep = ?

# 2. Define the model
Y1 = ?
Y1d = tf.nn.dropout(Y1, pkeep)

Y2 = tf.nn.relu(tf.matmul(?, W2) + B2)
Y2d = ?

Y3 = ?
Y3d = ?

Y4 = ?
Y4d = ?

Ylogits = ?
Y = ?

# 3. In the training step - provide the appropriate pkeep
# 4. In the training step, if you used a placeholder - adjust learning rate - according to exponential decay rate
def training_step(...):
    ...
    sess.run(?)
    ...
```

Task 3

The analysis tasks involve:

- Explain during grading the motivation behind learning rate decay.
- Explain during grading why dropout can be an effective regularization technique.
- For each of the programming tasks plot accuracy and loss, and analyze whether your additions influence the accuracy/loss and if yes, in what way.

4 Convolutional Neural Network

In this section you will change your neural network to use convolutional layers. In the previous tasks, we have reshaped our input from a matrix to a vector from the very beginning, however, this means that we lose precious information that could be including in the learning task. By getting rid of the matrix in the beginning we lose locality information. We lose the fact that the digits have lines and curves in them. The convolutional layers that we will add to the network will let us extract and compose higher level features from this locality information.

For this task you should read [this part](#) of the tensorflow tutorial in order to see how to setup a convolutional layer.

It is important to understand the shape of the input/output of your convolutional layers and what changes the shape of the input to that of the output. The weight variables from your convolutional layers should follow 4-dimensional tensors of *(patch_height, patch_width, input_channels, output_channels)*. This tensor defines the weights structure of one unit in the layer. The patch is the area of the image that the unit will try to look at. The input to the whole layer is a 4 dimensional tensor of *(batch_size, image_height, image_width, input_channels)*. The stride defines how to move the patch over the layer's input data in order to compute a unit's input data. Thus, the stride's structure is also 4 dimensional: *(batch_step, height_step, width_step, channel_step)*. Keep in mind that if the stride has any element bigger than 1 in its structure it will reduce that particular dimension in the output of the layer.

Next we provide a full layer parameter description.

Conv layer 1

The input data is $(batch_size, image_height, image_width, input_channels) = (100, 28, 28, 1)$

- stride of $(1, 1, 1, 1)$
- patch of 5×5
- input depth/channels of 1
- output depth/channels of 4

Conv layer 2

- stride of (1,2,2,1)
- patch of 5X5
- input depth/channels of 4 - previous output
- output depth/channels of 8

Conv layer 3

- stride of (1,2,2,1)
- patch of 4X4
- input depth/channels of 8 - previous output
- output depth/channels of 12

Densely connected layer

The output from the third layer is a four dimensional tensor of *(batch, height, width, depth)*. The densely connected layer is a relu layer like the ones defined in previous tasks and takes as input two dimension tensor of *(batch, values)*. You will need to reshape the tensor from having *height X width X depth* in matrix form to having a vector of *height * width * depth* elements.

- input structure - determine it from the previous layer output
- output structure: vector of 200 elements

Readout layer

The readout layer is the same softmax layer from the previous tasks.

- input structure: a vector of 200 elements
- output structure: vector of 10 elements

Task 4

Your programming part of this task includes:

1. Setup the network layer with 3 conv layers, 1 relu layer and 1 softmax layer with a GradientDescentOptimizer.
2. Change the optimizer to the AdamOptimizer.
3. Add a learning decay to the network.

4. Add regularization through dropout.

As a note - for programming subtask 1, all the changes you need to make are within variable declaration/initialization and model setup and as an example you have listing 6.

Your analysis part of this task includes:

- Define the output structure of the convolutional layers based on the given stride.
- For each of the programming subtasks 2-4 point out the changes that happen to the accuracy and error and explain why your modifications caused those changes.

Listing 6: "Convolutional network abstractions"

```
#define weight variable for a convolutional layer
W_I = tf.Variable(tf.truncated_normal([<patch_height>, <patch_width>, <input_channels>, <output_channels>],
    stddev=0.1))
#define convolutional layer in model
Y_I = tf.nn.relu(tf.nn.conv2d(<input>, <weights>, strides=[<batch_step>, <height_step>, <width_step>,
    <channel_step>], padding='SAME') + <bias>)
```

Task 5

5 Improving Predictions with Hyperparameter Optimization

There are now many different hyperparameters in your deep neural network, from the number of layers to the learning rate, choice of optimizer, to values for Dropout. In this task, you will train your network with many different combinations of hyperparameters, with the goal of improving predictions on the test dataset.

Docker/Python

If you are running Docker/Python, you are free to use whatever methods you prefer. You could, for example, implement your own random search or gridsearch method or just pick a number of experiments, and use a bash script to launch different experiments with different combinations of hyperparameters. One framework you could use is [hyperopt](#).

Hops

We suggest you use TFLauncher by Hops to supply your hyperparameters in a python dictionary. Each experiment will be launched as a separate Spark mapper, so you when you start your Jupyter notebook, be sure to increase the number of executors to match the number of experiments you will run. You will need 1 executor for every combination of hyperparameters you want to try out. So, if you have 10 experiments, you will need to launch 10 mappers.

Listing 7: "Hops Hyperparameters Code Template"

```
def mnist(learning_rate, dropout):  
    import tensorflow as tf  
    import numpy as np  
    from hops import tensorboard  
    from hops import hdfs  
  
    # Hyperparameters  
    args_dict = {'learning_rate': [0.001, 0.0001, 0.00001], 'dropout': [0.5, 0.7]}  
    from hops import util  
  
    # Generate grid of defined hyper parameters  
    args_dict_grid = util.grid_params(args_dict)  
    print(args_dict_grid)  
  
    from hops import tflauncher  
    tb_hdfs_dir = tflauncher.launch(spark, mnist, args_dict_grid)  
  
    from hops import tensorboard  
    tensorboard.visualize(spark, tb_hdfs_dir)
```

6 Distributed Training using TensorFlowOnSpark (Bonus task: 25% extra)

Re-implement your solution using TensorFlowOnSpark. This will require re-writing your application. TensorFlowOnSpark uses Parameter servers to distribute tasks among workers. You will probably have to do this task on hops.site, which has support for TensorFlowOnSpark. In your solution, you should compare the performance (training time) for an increasing number of executors on both normal TensorFlow and TensorFlowOnSpark.

7 Custom Spark Estimator/Evaluator/Experiment for Hops Tflauncher (Bonus task: 50% extra)

In this task, you have to write a custom Spark Estimator, Evaluator, and Experiment to 'wrap' Hop's Tflauncher library. The goal is that your application will now look more like a Spark ML Estimator/Pipeline model and can use Spark's classes for Grid Search and Evaluation, such as ParamGridBuilder and Evaluation. For details in Spark/Scala, see [here](#), and for PySpark, see [here](#). If you are intending to do this task, talk to us, as (from the last Databricks reference), it will be significantly easier to do it when Spark 2.3 is released. Release candidates for Spark 2.3 are expected in mid-November, and we can setup a Hopsworks cluster on which you can work on this task, with Spark 2.3 installed (www.hops.site has only Spark 2.2). This task could make a good project, if you wish to do this as your main project in ID2223.